

High-Performance Parallel and Distributed Systems – COM00036H Y3878747

Parallelisation approach [35%]

Assumption: the parallelised ports required no manual optimisations of the code (e.g. loop fusion/unrolling and vector intrinsics such as Intel's AVX/SSE for the OpenMP port).

Describe any performance analysis undertaken prior to parallelisation [100 words] (5%)

I extended the application to use a timer function to measure timings of the main operations with a `--countseconds`, or `-s`, flag [Fig 1]. As it was limited to nanosecond resolution, I also used the GNU Profiler [Fig 4]. Both corroborated that the `poisson()` function took the longest. Intel Advisor gave finer-grained timing information for individual loops within the functions [Fig 6]. Counting cache misses via L1 and L2 performance counters available on Viking [Fig 2] and PAPI was enabled via a `--countmisses`, or `-m`, flag, showing high data misses for `compute_tentative_velocity()` and `update_velocity()` [Fig 3].

Describe the approach taken to parallelise the application with OpenMP [100 words] (5%)

If iterations of a loop were independent, I used ``for`` clause in its pragma. This employed the SPMD technique for worksharing a loop across threads i.e. thread-level parallelism. However, race conditions could now occur where multiple threads updated a shared value simultaneously. I could have used critical regions and barriers to enforce synchronisation, but to avoid deadlocking, I used ``reduction`` clauses. For perfect loops, I reduced load imbalance via the ``collapse`` clause. Vectorisation techniques exploiting data-level parallelism could have been used via ``simd`` clauses (where backward dependencies between iterations of loops didn't exist) with vector intrinsics, but would reduce portability.

Describe the approach taken to parallelise the application with MPI [150 words] (15%)

Describe the approach taken to parallelise the application with CUDA [100 words] (10%)

All 2D arrays were converted to 1D arrays for CUDA compatibility. I used vectorisation techniques to exploit data-level parallelism in areas where backward dependencies between iterations of loops didn't exist. Data parallelism existed at the thread block scale, whereas task parallelism existed at the kernel level between independent tasks. Synchronisation barriers were used between sequentially dependent kernels. Atomic operations to avoid race conditions was only used once, when the latency would be offset by the number of other computations within a loop. To ensure efficient utilisation of resources despite inevitable latency, at least two threads should concurrently run on core.

Validation [15%]

Provide a brief description of the validation process undertaken and provide a summary of the results [150 words] (15%)

I extended the application to allow a new parameter, `-validate` or `-i`, which prints "Validated!" or "Invalidated" upon verification of the code. Without the `-i` tag used, the application defaults to run the standard code to completion. The standard code still runs to completion with `-i` parameter, however the serial, original version of the code is also run. The values of `u` and `v` calculated from the default code are then compared against the values calculated via the serial code. If the difference is greater than a certain value (defined by the `MAX_ERR` macro in the `validated()` function in `main.c`), the defaulted code is deemed invalid. I extended the application by moving the `main()` function from `vortex.c` into a new `main.c`, and kept the computation shared by both serial and parallel code portions within the main function to allow for less branching if validating, and thus reduce effects on timings.

Performance Evaluation [25%]

Outline the results collection and analysis process [100 words] (5%)

Timing functions [Fig 2] and PAPI counters [Fig 3] were used to evaluate the ports. Visit was used to visualise the output.

Provide appropriate figures and analysis demonstrating the performance of your OpenMP port [150 words] (10%)

The timing function [Fig 2] and PAPI counters [Fig 3] show that there was a significant decrease in the amount of time taken. However, some functions took slightly longer as well. Serialisation was more efficient for the functions: . This could have been mitigated by using certain `#pragma omp parallel for if (iters > 100)`, but only in areas where reductions were not used.

Provide appropriate figures and analysis demonstrating the performance of your MPI port [150 words] (5%)

Scalsca could have been used to evaluate the MPI port.

Provide appropriate figures and analysis demonstrating the performance of your CUDA port [150 words] (5%)

Nvprof was used to analyse the CUDA port. overhead when accessing elements of an array...

Comparative Analysis [25%]

Provide appropriate figures and provide an analysis of the performance and portability of your three ports [150 words] (10%)

The CUDA port is not that portable. MPI is portable.

The OpenMP code is the most portable, as it targets multiple architecture types, such as x86 or ARM CPUs. However, my port does not target heterogenous architectures like CUDA port can. MPI can be used in message passing between multiple CPUs on different nodes in a cluster enabling wider computation, however there is a larger communication overhead on a greater scale.

The block size and grid size of the cuda port is applicable to GPUs

Discuss the differences between your ports, highlighting any advantages or disadvantages that may exist [150 words] (15%)

OpenMP targets multiple architecture types, such as x86 or ARM CPUs, but newer versions can be heterogenous. However, my port does not target heterogenous architectures like CUDA port can. MPI can be used in message passing between multiple CPUs on different nodes in a cluster enabling wider computation, however there is a larger communication overhead on a greater scale. Programming models that targets heterogenous are the most performant and energy efficient. MPI targets distributed systems. But accelerator programming means lots of reengineering.

Appendix:

```
Timing Performance Summary
=====
Allocate Arrays Time: 0.000216
Problem Set Up Time: 0.003097
Average Compute Tentative Velocity Time: 0.001223
Average Poisson Time: 0.061926
Average Compute RHS Time: 0.000231
Average Update Velocity Time: 0.000276
Average Apply Boundary Conditions Time: 0.000054
Average Checkpoint Write Time: 0.000000

Total Problem Time: 132.713945
=====
```

[Fig 1: Pre-parallelisation timing summary, using configuration where timestep is 0.1 and end time 1 with file I/O disabled]

PAPI Preset Events					
Name	Code	Avail	Deriv	Description (Note)	
PAPI_L1_DCM	0x80000000	Yes	No	Level 1 data cache misses	
PAPI_L1_ICM	0x80000001	Yes	No	Level 1 instruction cache misses	
PAPI_L2_DCM	0x80000002	Yes	Yes	Level 2 data cache misses	
PAPI_L2_ICM	0x80000003	Yes	No	Level 2 instruction cache misses	
PAPI_L3_DCM	0x80000004	No	No	Level 3 data cache misses	
PAPI_L3_ICM	0x80000005	No	No	Level 3 instruction cache misses	
PAPI_L1_TCM	0x80000006	Yes	Yes	Level 1 cache misses	
PAPI_L2_TCM	0x80000007	Yes	No	Level 2 cache misses	
PAPI_L3_TCM	0x80000008	Yes	No	Level 3 cache misses	

[Fig 2: Pre-parallelisation performance counters available on Viking that I used in vortex_papi.c]

```

Step      0, Time: 1.00000000e-01 (del_t: 1.00000000e-01), Residual: 1.80316675e+01
Step      1, Time: 2.00000000e-01 (del_t: 1.00000000e-01), Residual: 1.37953704e+03
Step      2, Time: 3.00000000e-01 (del_t: 1.00000000e-01), Residual: 6.15407455e+05
Step      3, Time: 4.00000000e-01 (del_t: 1.00000000e-01), Residual: 4.32084181e+09
Step      4, Time: 5.00000000e-01 (del_t: 1.00000000e-01), Residual: 1.86754878e+17
Step      5, Time: 6.00000000e-01 (del_t: 1.00000000e-01), Residual: 3.79000681e+32
Step      6, Time: 7.00000000e-01 (del_t: 1.00000000e-01), Residual: 1.48040990e+63
Step      7, Time: 8.00000000e-01 (del_t: 1.00000000e-01), Residual: inf
Step      8, Time: 9.00000000e-01 (del_t: 1.00000000e-01), Residual: -nan
Step      9, Time: 1.00000000e+00 (del_t: 1.00000000e-01), Residual: -nan
Step     10, Time: 1.10000000e+00 (del_t: 1.00000000e-01), Residual: -nan
Step     11, Time: 1.10000000e+00, Residual: -nan
Simulation complete.
Cache Miss Summary:

Allocate Array Events:
Level 1 Cache Misses:  Data - 570, Instruction - 418, Total 988
Level 2 Cache Misses:  Data - 339, Instruction - 38, Total 377

Problem Set Up Events:
Level 1 Cache Misses:  Data - 12971, Instruction - 214, Total 13185

```

```
Level 1 Cache Misses:  Data - 12971, Instruction - 214, Total 13185
Level 2 Cache Misses:  Data - 8602, Instruction - 102, Total 8704

Average Compute Tentative Velocity Events:
Level 1 Cache Misses:  Data - 54611, Instruction - 175, Total 54787
Level 2 Cache Misses:  Data - 51616, Instruction - 163, Total 51779

Average Compute RHS Events:
Level 1 Cache Misses:  Data - 26597, Instruction - 73, Total 26671
Level 2 Cache Misses:  Data - 25452, Instruction - 71, Total 25524

Average Update Velocity Events:
Level 1 Cache Misses:  Data - 52675, Instruction - 64, Total 52739
Level 2 Cache Misses:  Data - 50934, Instruction - 63, Total 50998

Average Apply Boundary Conditions Events:
Level 1 Cache Misses:  Data - 2947, Instruction - 80, Total 3028
Level 2 Cache Misses:  Data - 1972, Instruction - 68, Total 2040


Total Write Events:
Level 1 Cache Misses:  Data - 0, Instruction - 0, Total 0
Level 2 Cache Misses:  Data - 0, Instruction - 0, Total 0

Remaining Events:
Level 1 Cache Misses:  Data - 2950, Instruction - 7515, Total 10465
Level 2 Cache Misses:  Data - 2536, Instruction - 5958, Total 8494
```

[Fig 3: Pre-parallelisation results yielded via PAPI in vortex_papi.c]

```
Flat profile:
Each sample counts as 0.01 seconds.
%   cumulative   self           self         total
time  seconds    seconds   calls   Ts/call   Ts/call   name
97.14    126.91    126.91           0      0.00     0.00    poisson
 1.80    129.26     2.35           0      0.00     0.00    compute_tentative_velocity
 0.35    129.72     0.46           0      0.00     0.00    update_velocity
 0.34    130.16     0.44           0      0.00     0.00    compute_rhs
 0.31    130.57     0.41           0      0.00     0.00    set_timestep_interval
 0.06    130.65     0.08           0      0.00     0.00    apply_boundary_conditions
 0.00    130.65     0.00           6      0.00     0.00    alloc_2d_array
 0.00    130.65     0.00           6      0.00     0.00    free_2d_array
 0.00    130.65     0.00           1      0.00     0.00    alloc_2d_char_array
```

[Fig 4: Pre-parallelisation GNU Profiler results]



Top Time-Consuming Loops ?		
Loop	Self Time?	Total Time?
loop in poisson at vortex.c:125	137.410s	137.410s
loop in poisson at vortex.c:153	83.394s	83.394s
loop in compute tentative_velocity at vortex.c:21	3.481s	3.481s
loop in compute tentative_velocity at vortex.c:44	3.169s	3.169s
loop in compute_rhs at vortex.c:85	1.776s	1.776s

Recommendations ?		
Enable the use of approximate division instructions	loop in compute tentative_velocity at vortex.c:21	
Enable the use of approximate division instructions	loop in compute tentative_velocity at vortex.c:44	
Enable the use of approximate division instructions	loop in compute_rhs at vortex.c:85	
Use the Intel short vector math library for vector intrinsics	loop in set timestep_interval at vortex.c:221	
Use the Intel short vector math library for vector intrinsics	loop in set timestep_interval at vortex.c:215	


[Fig 6: Pre-parallelisation Intel Advisor results summary]

Function Call Sites and Loops		Performance Issues		CPU Time		Type
				Total Time	Self Time	
[loop in compute tentative_velocity at vortex.c:44]		1 Unoptimized floating point operation processing possible		3.445s	3.445s	Vectorized (Body)
[loop in compute tentative_velocity at vortex.c:21]		1 Unoptimized floating point operation processing possible		3.389s	3.389s	Vectorized (Body)
[loop in poisson at vortex.c:125]				137.660s	137.660s	Vectorized (Body)

[Fig 7: Pre-parallelisation Intel Advisor Vectorization and Results' vectorisations]

Function Call Sites and Loops		Performance Issues		CPU Time		Why No Vectorization?
				Total Time	Self Time	
[loop in compute_rhs at vortex.c:85]		1 Unoptimized floating point operation processing possible		1.728s	1.728s	Compiler lacks sufficient information to vector
[loop in update_velocity at vortex.c:185]		1 Unoptimized floating point operation processing possible		0.396s	0.396s	Compiler lacks sufficient information to vector
[loop in update_velocity at vortex.c:194]		1 Unoptimized floating point operation processing possible		0.368s	0.368s	Compiler lacks sufficient information to vector
[loop in set timestep_interval at vortex.c:215]		1 Scalar math function call(s) present		0.708s	0.460s	
[loop in set timestep_interval at vortex.c:221]		1 Scalar math function call(s) present		0.704s	0.448s	
[loop in poisson at vortex.c:111]		1 Misaligned loop code present		0.184s	0.184s	Compiler lacks sufficient information to vector
[loop in apply_boundary_conditions at boundary.c:35]		1 Misaligned loop code present		0.120s	0.120s	
[loop in poisson at vortex.c:122]		1 Data type conversions present		83.596s	0.012s	
[loop in poisson at vortex.c:153]				83.320s	83.320s	Compiler lacks sufficient information to vector

[Fig 8: Non-vectorised code with performance issues via Intel Advisor Vectorization and Results pre-parallelisation]



[Fig 9: Pre-parallelisation Intel Advisor CPU/Memory Roofline Insights' roofline model]

Function Call Sites and Loops	Performance Issues	CPU Time	
		Total Time	Self Time
f_start		0.110s	0.000s
[loop in main at vortex.c:261]		0.110s	0.000s
f_main		0.110s	0.000s
f_poisson		0.110s	0.000s
[loop in poisson at vortex.c:122]		0.110s	0.000s
[loop in poisson at vortex.c:125]	1 Data type conversions present	0.062s	0.062s
[loop in poisson at vortex.c:124]		0.062s	0.000s
[loop in poisson at vortex.c:124]		0.062s	0.000s
[loop in poisson at vortex.c:153]		0.048s	0.048s
[loop in poisson at vortex.c:152]		0.048s	0.000s

[Fig 10: Pre-parallelisation Intel Advisor Threading' survey]

Recommendations

Use the smallest data type loop in [poisson](#) at [vortex.c:122](#)

Top Time-Consuming Loops

Consider adding parallel site and task annotations around these time-consuming loops found during Survey analysis.

Loop	Self Time	Total Time	Trip Counts
loop in poisson at vortex.c:122	<0.001s	0.110s	100
loop in main at vortex.c:261	<0.001s	0.110s	1
loop in poisson at vortex.c:125	0.062s	0.062s	128
loop in poisson at vortex.c:124	<0.001s	0.062s	512
loop in poisson at vortex.c:124	<0.001s	0.062s	2

Function Call Sites and Loops	Trip Counts		Performance Issues	CPU Time	
	Average	Call Count		Total Time	Self Time
f_start				0.110s	0.000s
f_main				0.110s	0.000s
f_poisson				0.110s	0.000s
[loop in poisson at vortex.c:122]	100	1		0.110s	0.000s
[loop in main at vortex.c:261]	1	1		0.110s	0.000s
[loop in poisson at vortex.c:124]	2	100		0.062s	0.000s
[loop in poisson at vortex.c:124]	512	200		0.062s	0.000s
[loop in poisson at vortex.c:125]	128	102400		0.062s	0.062s
[loop in poisson at vortex.c:152]	512	100		0.048s	0.000s
[loop in poisson at vortex.c:153]	128	51200		0.048s	0.048s

Loop at 122 is called once but has iteration of 100

```

Timing Performance Summary
=====
Allocate Arrays Time: 0.000193
Problem Set Up Time: 0.002893
Average Compute Tentative Velocity Time: 0.001325
Average Poisson Time: 0.049781
Average Compute RHS Time: 0.000250
Average Update Velocity Time: 0.000300
Average Apply Boundary Conditions Time: 0.000058
Average Checkpoint Write Time: 0.000000

Total Problem Time: 67.801544
=====

```

[Timing summary using OpenMP port, using default configuration with file I/O disabled]

```

=====
Cache Miss Summary=====

Allocate Array Events:
Level 1 Cache Misses:   Data - 717, Instruction - 467, Total 1184
Level 2 Cache Misses:   Data - 368, Instruction - 58, Total 426

Problem Set Up Events:
Level 1 Cache Misses:   Data - 14469, Instruction - 261, Total 14730
Level 2 Cache Misses:   Data - 9360, Instruction - 158, Total 9518

Average Compute Tentative Velocity Events:
Level 1 Cache Misses:   Data - 53939, Instruction - 199, Total 54138
Level 2 Cache Misses:   Data - 50437, Instruction - 178, Total 50615

Average Compute RHS Events:
Level 1 Cache Misses:   Data - 26336, Instruction - 106, Total 26442
Level 2 Cache Misses:   Data - 25405, Instruction - 104, Total 25510

```



```
Average Update Velocity Events:
Level 1 Cache Misses:  Data - 52698, Instruction - 108,
Total 52806
Level 2 Cache Misses:  Data - 50650, Instruction - 105,
Total 50756

Average Apply Boundary Conditions Events:
Level 1 Cache Misses:  Data - 2956, Instruction - 54, To
tal 3011
Level 2 Cache Misses:  Data - 1919, Instruction - 46, To
tal 1966

Average Checkpoint Write Events:
Level 1 Cache Misses:  Data - 0, Instruction - 0, Total
0
Level 2 Cache Misses:  Data - 0, Instruction - 0, Total
0

Remaining Looping Events:
Level 1 Cache Misses:  Data - 1944, Instruction - 4880,
Total 6824
Level 2 Cache Misses:  Data - 1690, Instruction - 3842,
Total 5532
=====
```

[Cache miss summary using OpenMP port, using configuration where timestep is 0.1 and end time 1 with file I/O disabled]

[Intel Advisor CPU/Memory Roofline Insights]

[Threading Survey and Roofline results]